

DATA SCIENCE SEMESTER – V			
Course Code	M23BCI503	CIE Marks	50
Number of Lecture Hours/Week(L: T: P: S)	(4:0:0:0)	SEE Marks	50
Total Number of Lecture Hours	40 hours Theory	Total Marks	100
Credits	04	Exam Hours	03
Course Objectives: <ol style="list-style-type: none"> 1. Introduce fundamental concepts and techniques in data science. 2. Equip students with practical skills for data analysis, visualization, and machine learning. 3. Foster an understanding of real-world applications of data science, such as in recommendation systems and predictive modelling. 4. Prepare students for advanced study or careers in data science and related fields. 			
Module -1			
Introduction to Data Science Introduction: What is Data Science? Big Data and Data Science hype – and getting past the hype, Why now? Datafication, Current landscape of perspectives, Skill sets needed. Statistical Inference: Populations and samples, Statistical modelling, probability distributions, fitting a model.			
Module -2			
Exploratory Data Analysis and the Data Science Process Basic tools (plots, graphs and summary statistics) of EDA, Philosophy of EDA, The Data Science Process, Case Study: Real Direct (online real estate firm). Three Basic Machine Learning Algorithms: Linear Regression, k-Nearest Neighbours (k- NN), k-means.			
Module -3			
One More Machine Learning Algorithm and Usage in Applications - Motivating application: Filtering Spam - Why Linear Regression and k-NN are poor choices for Filtering Spam - Naive Bayes and why it works for Filtering Spam - Data Wrangling: APIs and other tools for scrapping the Web.			
Module -4			
Feature Generation and Feature Selection (Extracting Meaning From Data) - Motivating application: user (customer) retention - Feature Generation (brainstorming, role of domain expertise, and place for imagination) - Feature Selection algorithms – Filters; Wrappers; Decision Trees; Random Forests.			
Module -5			
Recommendation Systems: Building a User-Facing Data Product - Algorithmic ingredients of a Recommendation Engine - Dimensionality Reduction - Singular Value Decomposition - Principal Component Analysis - Exercise: build your own recommendation system Mining Social-Network Graphs - Social networks as graphs - Clustering of graphs - Direct discovery of communities in graphs - Partitioning of graphs - Neighbourhood properties in graphs.			
TEXTBOOKS: <ol style="list-style-type: none"> 1. Cathy O’Neil and Rachel Schutt. Doing Data Science, Straight Talk FromThe Frontline. O’Reilly. 2014. 2. Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost and Tom Fawcett, O'Reilly Media, 2013. 			
REFERENCE BOOKS: <ol style="list-style-type: none"> 1. Jure Leskovek, Anand Rajaraman and Jeffrey Ullman. Mining of Massive Datasets. v2.1, Cambridge University Press. 2014. (free online). 2. Kevin P. Murphy. Machine Learning: A Probabilistic Perspective. ISBN 0262018020. 2013. 3. Python for Data Analysis by Wes McKinney, O'Reilly Media, 2017. 			